

Customer Churn Prediction and Analysis

This project aims to predict which customers are more likely to discontinue their subscription service, also known as churn prediction. Using the Telco Customer Churn dataset from Kaggle, I will apply data analysis techniques to classify customers as "returned" or "churned". The dataset includes information regarding the customers' demographics, the services they subscribed to, whether the customer left within the last month, and account information. The goal is to compare simple machine learning models with a deep learning approach to see which provides better predictions and business insights.

```
In [20]: # import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, f1_score, roc_auc_score, classification_report
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.callbacks import EarlyStopping
```

```
In [2]: # 1. Loading the Dataset
df = pd.read_csv('dataset/Telco-Customer-Churn.csv')
df.head()
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 100 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   rev_Mean          99643 non-null   float64
 1   mou_Mean          99643 non-null   float64
 2   totmrc_Mean       99643 non-null   float64
 3   da_Mean           99643 non-null   float64
 4   ovr mou_Mean      99643 non-null   float64
 5   ovr rev_Mean      99643 non-null   float64
 6   vce ovr_Mean      99643 non-null   float64
 7   dat ovr_Mean      99643 non-null   float64
 8   roam_Mean         99643 non-null   float64
 9   change_mou        99109 non-null   float64
 10  change_rev        99109 non-null   float64
 11  drop_vce_Mean     100000 non-null   float64
 12  drop_dat_Mean     100000 non-null   float64
 13  blck_vce_Mean     100000 non-null   float64
 14  blck_dat_Mean     100000 non-null   float64
 15  unan_vce_Mean     100000 non-null   float64
 16  unan_dat_Mean     100000 non-null   float64
 17  plcd_vce_Mean     100000 non-null   float64
 18  plcd_dat_Mean     100000 non-null   float64
 19  recv_vce_Mean     100000 non-null   float64
 20  recv_sms_Mean     100000 non-null   float64
 21  comp_vce_Mean     100000 non-null   float64
 22  comp_dat_Mean     100000 non-null   float64
 23  custcare_Mean     100000 non-null   float64
 24  ccrndmou_Mean     100000 non-null   float64
 25  cc_mou_Mean       100000 non-null   float64
 26  inonemin_Mean     100000 non-null   float64
 27  three way_Mean    100000 non-null   float64
 28  mou_cvce_Mean     100000 non-null   float64
 29  mou_cdat_Mean     100000 non-null   float64
 30  mou_rvce_Mean     100000 non-null   float64
 31  owylis_vce_Mean   100000 non-null   float64
 32  mouowylisv_Mean   100000 non-null   float64
 33  iwy lis_vce_Mean 100000 non-null   float64
 34  mou iwy lisv_Mean 100000 non-null   float64
 35  peak_vce_Mean     100000 non-null   float64
 36  peak_dat_Mean     100000 non-null   float64
 37  mou_peav_Mean     100000 non-null   float64
 38  mou_ pead_Mean    100000 non-null   float64
 39  opk_vce_Mean      100000 non-null   float64
 40  opk_dat_Mean      100000 non-null   float64
 41  mou_ opkv_Mean    100000 non-null   float64
 42  mou_ opkd_Mean    100000 non-null   float64
 43  drop_blk_Mean     100000 non-null   float64
 44  attempt_Mean      100000 non-null   float64
 45  complete_Mean     100000 non-null   float64
 46  call fwdv_Mean    100000 non-null   float64
 47  call wait_Mean    100000 non-null   float64
 48  churn             100000 non-null   int64 
 49  months            100000 non-null   int64 
 50  uniqsubs          100000 non-null   int64
```

```
51 actvsubs          100000 non-null int64
52 new_cell          100000 non-null object
53 crclscod          100000 non-null object
54 asl_flag           100000 non-null object
55 totcalls          100000 non-null int64
56 totmou            100000 non-null float64
57 totrev             100000 non-null float64
58 adjrev             100000 non-null float64
59 adjmou             100000 non-null float64
60 adjqty              100000 non-null int64
61 avgrev             100000 non-null float64
62 avgmou             100000 non-null float64
63 avgqty              100000 non-null float64
64 avg3mou            100000 non-null int64
65 avg3qty             100000 non-null int64
66 avg3rev             100000 non-null int64
67 avg6mou            97161 non-null float64
68 avg6qty             97161 non-null float64
69 avg6rev             97161 non-null float64
70 prizm_social_one   92612 non-null object
71 area                99960 non-null object
72 dualband            99999 non-null object
73 refurb_new           99999 non-null object
74 hnd_price            99153 non-null float64
75 phones               99999 non-null float64
76 models               99999 non-null float64
77 hnd_webcap           89811 non-null object
78 truck                98268 non-null float64
79 rv                   98268 non-null float64
80 ownrent              66294 non-null object
81 lor                  69810 non-null float64
82 dwlltype              68091 non-null object
83 marital               98268 non-null object
84 adults                 76981 non-null float64
85 infobase              77921 non-null object
86 income                 74564 non-null float64
87 numbcars              50634 non-null float64
88 HHstatin              62077 non-null object
89 dwllsize              61692 non-null object
90 forgntvl              98268 non-null float64
91 ethnic                 98268 non-null object
92 kid0_2                 98268 non-null object
93 kid3_5                 98268 non-null object
94 kid6_10                 98268 non-null object
95 kid11_15                 98268 non-null object
96 kid16_17                 98268 non-null object
97 creditcd                 98268 non-null object
98 eqpdays                 99999 non-null float64
99 Customer_ID            100000 non-null int64
dtypes: float64(69), int64(10), object(21)
memory usage: 76.3+ MB
```

```
In [3]: # 2. Data preprocessing
df["churn"] = LabelEncoder().fit_transform(df["churn"])

# One-hot coding
```

```

df = pd.get_dummies(df, drop_first=True)

# Separate features and target
X = df.drop("churn", axis=1)
y = df["churn"]

# Create an imputer to replaces NaN with the column mean
from sklearn.impute import SimpleImputer

imputer = SimpleImputer(strategy='mean')
X = pd.DataFrame(imputer.fit_transform(X), columns=X.columns)

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Features scaling
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

print("Data preprocessing complete.")
print(f"Training set shape: {X_train.shape}")
print(f"Test set shape: {X_test.shape}")
print(f"NaN values? {pd.isna(X_train).any()}")

```

Data preprocessing complete.
Training set shape: (80000, 210)
Test set shape: (20000, 210)
NaN values? False

In [22]:

```

# 2.1. Data exploration
# Check dataset dimensions and types
print("Dataset shape:", df.shape)
print("\nData types:")
print(df.dtypes.head(10))

# Look at missing values
print("\nMissing values per column:")
print(df.isnull().sum().sort_values(ascending=False).head(15))

# Churn distribution
print(f"Churn Rate: {y.mean():.2%}")
plt.figure(figsize=(5,4))
sns.countplot(x="churn", data=df, palette="coolwarm")
plt.title("Churn Distribution (0 = Retained, 1 = Churned)")
plt.show()

```

```
Dataset shape: (100000, 211)
```

```
Data types:
```

```
rev_Mean      float64
mou_Mean      float64
totmrc_Mean   float64
da_Mean       float64
ovrmou_Mean   float64
ovrrev_Mean   float64
vceovr_Mean   float64
datovr_Mean   float64
roam_Mean     float64
change_mou    float64
dtype: object
```

```
Missing values per column:
```


```
numbcars      49366
lor           30190
income         25436
adults         23019
avg6rev       2839
avg6qty       2839
avg6mou       2839
rv             1732
forgntvl      1732
truck          1732
change_mou    891
change_rev    891
hnd_price     847
totmrc_Mean   357
ovrmou_Mean   357
dtype: int64
```

```
Churn Rate: 49.56%
```

```
C:\Users\anhth\AppData\Local\Temp\ipykernel_15504\3314583481.py:14: FutureWarning:
C:\Users\anhth\AppData\Local\Temp\ipykernel_22600\3314583481.py:14: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.1
4.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.
```

```
sns.countplot(x="churn", data=df, palette="coolwarm")
```



```
In [15]: # 3. Baseline models
# Logistic Regression
lr = LogisticRegression(max_iter=1000, n_jobs=-1) #Use all CPU cores for parallel t
lr.fit(X_train, y_train)
lr_preds = lr.predict(X_test)

# Random Forest
rf = RandomForestClassifier(
    n_estimators=100,
    random_state=42,
    n_jobs=-1,      # use all CPU cores
    max_depth = 10) # limit tree depth to prevent overfitting
rf.fit(X_train, y_train)
rf_preds = rf.predict(X_test)

# 3.1. Model evaluation
for name, preds in [('Logistic Regression', lr_preds), ('Random Forest', rf_preds)]:
    print(f'{name}')
    print(f"Accuracy: {accuracy_score(y_test, preds):.3f}")
    print(f"F1 Score: {f1_score(y_test, preds):.3f}")
    print(f"ROC AUC: {roc_auc_score(y_test, preds):.3f}\n")

# 3.2. Compare baseline models results
models_list = ['Logistic Regression', 'Random Forest']
metrics = {
    'Accuracy': [0.592, 0.617],
    'F1 Score': [0.590, 0.628],
    'ROC AUC': [0.592, 0.617]
}

x = np.arange(len(models_list))
width = 0.25
```

```

fig, ax = plt.subplots(figsize=(12, 6))
for i, (metric, values) in enumerate(metrics.items()):
    ax.bar(x + i*width, values, width, label=metric)

ax.set_xlabel('Models')
ax.set_ylabel('Score')
ax.set_title('Model Performance Comparison')
ax.set_xticks(x + width)
ax.set_xticklabels(models_list)
ax.legend()
ax.set_xlim(0.5, 0.7)
plt.tight_layout()
plt.show()

```

Logistic Regression

Accuracy: 0.592

F1 Score: 0.590


ROC AUC: 0.592

Random Forest

Accuracy: 0.617

F1 Score: 0.628

ROC AUC: 0.617



In [6]: # 3.3. Feature importance from Random Forest

```


importances = rf.feature_importances_
feature_names = X.columns
importance_df = pd.DataFrame({
    'feature': feature_names,
    'importance': importances
}).sort_values('importance', ascending=False).head(10)

```

```

plt.figure(figsize=(10, 6))
plt.barh(importance_df['feature'], importance_df['importance'])
plt.xlabel('Importance')
plt.title('Top 10 Most Important Features for Churn Prediction')
plt.gca().invert_yaxis()
plt.show()

```



In [7]: # 4. Deep Neural Network model

```

model = models.Sequential([
    # Input Layer with more neurons for complex patterns
    layers.Input(shape=(X_train.shape[1],)),
    layers.Dense(128, activation='relu'),
    layers.BatchNormalization(),
    layers.Dropout(0.3),

    # hidden Layer
    layers.Dense(64, activation='relu'),
    layers.BatchNormalization(),
    layers.Dropout(0.3),
    layers.Dense(32, activation='relu'),
    layers.BatchNormalization(),
    layers.Dropout(0.2),

    # Output Layer
    layers.Dense(1, activation='sigmoid')
])

model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['accuracy', tf.keras.metrics.AUC(name='auc')])

# add EarlyStopping for better training
callbacks = [

```

```
# Stop training if validation loss doesn't improve for 5 epochs
EarlyStopping(
    monitor='val_loss',
    patience=5,
    restore_best_weights=True,
    verbose=1
)
]

# training
history = model.fit(
    X_train, y_train,
    epochs=30,
    batch_size=128,
    validation_split=0.2,
    callbacks=callbacks,
    verbose=1
)

loss, accuracy, auc = model.evaluate(X_test, y_test)
print(f"Test Accuracy: {accuracy:.3f}, Test AUC: {auc:.3f}")
```

Epoch 1/30
500/500 8s 7ms/step - accuracy: 0.5294 - auc: 0.5399 - loss: 0.
7317 - val_accuracy: 0.5756 - val_auc: 0.6037 - val_loss: 0.6777

Epoch 2/30
500/500 3s 6ms/step - accuracy: 0.5652 - auc: 0.5899 - loss: 0.
6831 - val_accuracy: 0.5897 - val_auc: 0.6222 - val_loss: 0.6707

Epoch 3/30
500/500 3s 6ms/step - accuracy: 0.5841 - auc: 0.6165 - loss: 0.
6721 - val_accuracy: 0.5938 - val_auc: 0.6285 - val_loss: 0.6683

Epoch 4/30
500/500 3s 6ms/step - accuracy: 0.5926 - auc: 0.6284 - loss: 0.
6676 - val_accuracy: 0.6002 - val_auc: 0.6370 - val_loss: 0.6643

Epoch 5/30
500/500 3s 6ms/step - accuracy: 0.5983 - auc: 0.6358 - loss: 0.
6641 - val_accuracy: 0.6036 - val_auc: 0.6421 - val_loss: 0.6629

Epoch 6/30
500/500 3s 6ms/step - accuracy: 0.6047 - auc: 0.6430 - loss: 0.
6612 - val_accuracy: 0.6060 - val_auc: 0.6444 - val_loss: 0.6616

Epoch 7/30
500/500 3s 6ms/step - accuracy: 0.6069 - auc: 0.6489 - loss: 0.
6582 - val_accuracy: 0.6079 - val_auc: 0.6496 - val_loss: 0.6590

Epoch 8/30
500/500 3s 6ms/step - accuracy: 0.6101 - auc: 0.6529 - loss: 0.
6562 - val_accuracy: 0.6097 - val_auc: 0.6515 - val_loss: 0.6586

Epoch 9/30
500/500 3s 6ms/step - accuracy: 0.6169 - auc: 0.6592 - loss: 0.
6531 - val_accuracy: 0.6106 - val_auc: 0.6530 - val_loss: 0.6564

Epoch 10/30
500/500 3s 6ms/step - accuracy: 0.6178 - auc: 0.6622 - loss: 0.
6517 - val_accuracy: 0.6142 - val_auc: 0.6553 - val_loss: 0.6561

Epoch 11/30
500/500 3s 6ms/step - accuracy: 0.6200 - auc: 0.6659 - loss: 0.
6492 - val_accuracy: 0.6140 - val_auc: 0.6553 - val_loss: 0.6560

Epoch 12/30
500/500 3s 6ms/step - accuracy: 0.6230 - auc: 0.6705 - loss: 0.
6469 - val_accuracy: 0.6108 - val_auc: 0.6533 - val_loss: 0.6565

Epoch 13/30
500/500 3s 6ms/step - accuracy: 0.6260 - auc: 0.6743 - loss: 0.
6445 - val_accuracy: 0.6144 - val_auc: 0.6574 - val_loss: 0.6548

Epoch 14/30
500/500 3s 6ms/step - accuracy: 0.6261 - auc: 0.6746 - loss: 0.
6446 - val_accuracy: 0.6139 - val_auc: 0.6577 - val_loss: 0.6546

Epoch 15/30
500/500 3s 6ms/step - accuracy: 0.6293 - auc: 0.6789 - loss: 0.
6419 - val_accuracy: 0.6131 - val_auc: 0.6586 - val_loss: 0.6547

Epoch 16/30
500/500 3s 6ms/step - accuracy: 0.6338 - auc: 0.6847 - loss: 0.
6383 - val_accuracy: 0.6139 - val_auc: 0.6536 - val_loss: 0.6568

Epoch 17/30
500/500 3s 6ms/step - accuracy: 0.6356 - auc: 0.6874 - loss: 0.
6366 - val_accuracy: 0.6139 - val_auc: 0.6563 - val_loss: 0.6551

Epoch 18/30
500/500 3s 6ms/step - accuracy: 0.6360 - auc: 0.6892 - loss: 0.
6354 - val_accuracy: 0.6112 - val_auc: 0.6550 - val_loss: 0.6561

Epoch 19/30
500/500 3s 6ms/step - accuracy: 0.6386 - auc: 0.6923 - loss: 0.


```
6334 - val_accuracy: 0.6161 - val_auc: 0.6576 - val_loss: 0.6560
Epoch 19: early stopping
Restoring model weights from the end of the best epoch: 14.
625/625 2s 3ms/step - accuracy: 0.6090 - auc: 0.6527 - loss: 0.6569
Test Accuracy: 0.609, Test AUC: 0.653
```

```
In [8]: # 4.1. Model Evaluation
loss, accuracy, auc = model.evaluate(X_test, y_test, verbose=0)
print(f"\nTest Results:")
print(f" Loss: {loss:.4f}")
print(f" Accuracy: {accuracy:.3f}")
print(f" ROC AUC: {auc:.3f}")

# Visualization
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend()
plt.title('DNN Accuracy Over Epochs')
plt.show()
```

Test Results:

```
Loss: 0.6569
Accuracy: 0.609
ROC AUC: 0.653
```



```
In [9]: # 4.2. Evaluate Recall and Precision of DNN
dnn_preds = (model.predict(X_test, verbose=0) > 0.5).astype(int).flatten()
# Get predictions
dnn_preds = (model.predict(X_test, verbose=0) > 0.5).astype(int).flatten()

print("DNN Classification Report")
print("*"*60)
print(classification_report(y_test, dnn_preds))
```

	precision	recall	f1-score	support
0	0.61	0.62	0.62	10021
1	0.61	0.59	0.60	9979
0	0.62	0.60	0.61	10021
1	0.61	0.63	0.62	9979
accuracy			0.61	20000
macro avg	0.61	0.61	0.61	20000
weighted avg	0.61	0.61	0.61	20000

Deep Neural Network Architecture Tuning

After exploring the basic DNN architecture, I want to test different network architectures to find the optimal configuration:

- Varying number of layers
- Different layer sizes
- Different dropout rates

```
In [10]: # 5.1. DNN Architecture Tuning
print("Testing different DNN architectures...\n")

tuning_results = []

# Define architectures to test
architectures = [
    {'name': 'Baseline', 'layers': [128, 64, 32], 'dropout': 0.3},
    {'name': 'Deeper', 'layers': [256, 128, 64, 32], 'dropout': 0.3},
    {'name': 'Wider', 'layers': [256, 128, 64], 'dropout': 0.4},
    {'name': 'Simpler', 'layers': [128, 64], 'dropout': 0.2},
]

for arch in architectures:
    print(f"Training {arch['name']} architecture: {arch['layers']}")

    # Build model
    model_test = models.Sequential()
    model_test.add(layers.Input(shape=(X_train.shape[1],)))

    # Add Layers
```

```

for units in arch['layers']:
    model_test.add(layers.Dense(units, activation='relu'))
    model_test.add(layers.BatchNormalization())
    model_test.add(layers.Dropout(arch['dropout']))

# Output layer
model_test.add(layers.Dense(1, activation='sigmoid'))

# Compile
model_test.compile(
    optimizer='adam',
    loss='binary_crossentropy',
    metrics=['accuracy', tf.keras.metrics.AUC(name='auc')])
)

# Train with early stopping
history = model_test.fit(
    X_train, y_train,
    epochs=20,
    batch_size=128,
    validation_split=0.2,
    callbacks=[EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True),
    verbose=0
)

# Evaluate
loss, acc, auc = model_test.evaluate(X_test, y_test, verbose=0)

tuning_results.append({
    'name': arch['name'],
    'architecture': str(arch['layers']),
    'dropout': arch['dropout'],
    'accuracy': acc,
    'auc': auc
})

print(f" Test Accuracy: {acc:.3f}, Test AUC: {auc:.3f}\n")

# Display results
print("\n" + "="*60)
print("DNN Architecture Tuning Results")
print("="*60)
tuning_df = pd.DataFrame(tuning_results)
print(tuning_df.to_string(index=False))

# Find best architecture
best_arch = tuning_df.loc[tuning_df['auc'].idxmax()]
print(f"\nBest Architecture: {best_arch['name']} ")
print(f" Architecture: {best_arch['architecture']} ")
print(f" Dropout: {best_arch['dropout']} ")
print(f" Accuracy: {best_arch['accuracy']:.3f} ")
print(f" ROC AUC: {best_arch['auc']:.3f} ")

# 5.2. Visualize tuning results
print(f" ROC AUC: {best_arch['auc']:.3f}")

```

Testing different DNN architectures...

Training Baseline architecture: [128, 64, 32]

Test Accuracy: 0.614, Test AUC: 0.657

Training Deeper architecture: [256, 128, 64, 32]

Test Accuracy: 0.610, Test AUC: 0.653

Training Wider architecture: [256, 128, 64]

Test Accuracy: 0.615, Test AUC: 0.656

Training Simpler architecture: [128, 64]

Test Accuracy: 0.609, Test AUC: 0.650

=====

DNN Architecture Tuning Results

=====

name	architecture	dropout	accuracy	auc
Baseline	[128, 64, 32]	0.3	0.61355	0.657162
Deeper	[256, 128, 64, 32]	0.3	0.61020	0.653437
Wider	[256, 128, 64]	0.4	0.61530	0.655563
Simpler	[128, 64]	0.2	0.60870	0.649687


Best Architecture: Baseline

Architecture: [128, 64, 32]

Dropout: 0.3

Accuracy: 0.614

ROC AUC: 0.657



```
In [13]: # 4.2. Visualize tuning results
fig, ax = plt.subplots(figsize=(10, 6))

# Bar chart of architectures
```


```

arch_names = [r['name'] for r in tuning_results]
accuracies = [r['accuracy'] for r in tuning_results]
aucs = [r['auc'] for r in tuning_results]

x = np.arange(len(arch_names))
width = 0.35

ax.bar(x - width/2, accuracies, width, label='Accuracy', color='steelblue')
ax.bar(x + width/2, aucs, width, label='ROC AUC', color='coral')
ax.set_xlabel('Architecture')
ax.set_ylabel('Score')
ax.set_title('DNN Architecture Comparison')
ax.set_xticks(x)
ax.set_xticklabels(arch_names, rotation=45)
ax.legend()
ax.set_xlim(0.5, 0.7)
ax.grid(axis='y', alpha=0.3)
plt.tight_layout()
plt.show()
ax.grid(axis='y', alpha=0.3)

```



Comparing all models: DNN, Logistic Regression, Random Forest

This section evaluates and compares the performance of DNN against LR and RF.

```

In [16]: # 6. Compare DNN with baseline models
print("Models Comparison")
print("*"*60)

```

```

# 6.1. Overall performance comparison
# Get predictions for fair comparison
dnn_preds = (model.predict(X_test, verbose=0) > 0.5).astype(int).flatten()

print(f"\nLogistic Regression:")
print(f" Accuracy: {accuracy_score(y_test, lr_preds):.3f}")
print(f" F1 Score: {f1_score(y_test, lr_preds):.3f}")
print(f" ROC AUC: {roc_auc_score(y_test, lr_preds):.3f}")

print(f"\nRandom Forest:")
print(f" Accuracy: {accuracy_score(y_test, rf_preds):.3f}")
print(f" F1 Score: {f1_score(y_test, rf_preds):.3f}")
print(f" ROC AUC: {roc_auc_score(y_test, rf_preds):.3f}")

print(f"\nDeep Neural Network:")
print(f" Accuracy: {accuracy_score(y_test, dnn_preds):.3f}")
print(f" F1 Score: {f1_score(y_test, dnn_preds):.3f}")
print(f" ROC AUC: {auc:.3f}")

models_list = ['Logistic Regression', 'Random Forest', 'Deep Neural Net']
metrics = {
    'Accuracy': [0.592, 0.617, 0.610],
    'F1 Score': [0.590, 0.628, 0.624],
    'ROC AUC': [0.592, 0.617, 0.653]
}

x = np.arange(len(models_list))
width = 0.25

fig, ax = plt.subplots(figsize=(12, 6))
for i, (metric, values) in enumerate(metrics.items()):
    ax.bar(x + i*width, values, width, label=metric)

ax.set_xlabel('Models')
ax.set_ylabel('Score')
ax.set_title('Model Performance Comparison')
ax.set_xticks(x + width)
ax.set_xticklabels(models_list)
ax.legend()
ax.set_ylim(0.5, 0.7)
plt.tight_layout()
plt.show()

```

Models Comparison

Logistic Regression:


Accuracy: 0.592
F1 Score: 0.590
ROC AUC: 0.592

Random Forest:

Accuracy: 0.617
F1 Score: 0.628
ROC AUC: 0.617

Deep Neural Network:

Accuracy: 0.609
F1 Score: 0.603
ROC AUC: 0.650
Accuracy: 0.611
F1 Score: 0.617
ROC AUC: 0.653



```
In [12]: # 6.2. Individual customer prediction breakdown to see how each model produces a pr
# generate sample customer data
customer_idx = 0
customer_features = X_test[customer_idx]
actual_churn = y_test.iloc[customer_idx]

print("*"*80)
print("CUSTOMER PREDICTION BREAKDOWN")
print("*"*80)

print(f"\n INPUT FEATURES")
print(f" Number of features: {len(customer_features)}")
print(f" Sample values: {customer_features[:10]}")

print(f"\n ACTUAL OUTCOME:")
print(f" churn = {actual_churn} ({'CHURNED' if actual_churn == 1 else 'STAYED'})")
```

```

print(f"\n MODEL PREDICTIONS:")

# Logistic Regression
lr_pred = lr.predict([customer_features])[0]
lr_prob = lr.predict_proba([customer_features])[0]
print(f"  Logistic Regression:")
print(f"    Prediction: {lr_pred} ({'CHURN' if lr_pred == 1 else 'STAY'})")
print(f"    Probabilities: [Stay: {lr_prob[0]:.1%}, Churn: {lr_prob[1]:.1%}]")

# Random Forest
rf_pred = rf.predict([customer_features])[0]
rf_prob = rf.predict_proba([customer_features])[0]
print(f"  Random Forest:")
print(f"    Prediction: {rf_pred} ({'CHURN' if rf_pred == 1 else 'STAY'})")
print(f"    Probabilities: [Stay: {rf_prob[0]:.1%}, Churn: {rf_prob[1]:.1%}]")

# Deep Neural Network
dnn_prob = model.predict(np.array([customer_features]), verbose=0)[0][0]
dnn_pred = 1 if dnn_prob > 0.5 else 0
print(f"  Deep Neural Network:")
print(f"    Prediction: {'Churn' if dnn_pred == 1 else 'Stay'}")
print(f"    Churn Probability: {dnn_prob:.1%}")

print("=*80")

```

=====

CUSTOMER PREDICTION BREAKDOWN

=====

INPUT FEATURES

```

Number of features: 210
Sample values: [-0.32711824  0.34981012 -0.05048493 -0.40527053 -0.17198534 -0.164
82292
-0.15809594 -0.08688812 -0.08107395  0.11658782]
```

ACTUAL OUTCOME:

```
churn = 1 (CHURNED)
```

MODEL PREDICTIONS:

Logistic Regression:

```
Prediction: 1 (CHURN)
```

```
Probabilities: [Stay: 44.1%, Churn: 55.9%]
```

Random Forest:

```
Prediction: 1 (CHURN)
```

```
Probabilities: [Stay: 39.5%, Churn: 60.5%]
```

Deep Neural Network:

```
Prediction: Stay
```

```
Churn Probability: 41.6%
```

=====

In []: # 6.2. View sample predictions from all models

```
# Get predictions and churn probabilities from all models
```

```
lr_preds = lr.predict(X_test)
```

```
lr_probs = lr.predict_proba(X_test)[:, 1]
```

```
rf_preds = rf.predict(X_test)
```

```

rf_probs = rf.predict_proba(X_test)[:, 1]

dnn_probs = model.predict(X_test, verbose=0).flatten()
dnn_preds = (dnn_probs > 0.5).astype(int)

# Create comparison DataFrame
predictions_df = pd.DataFrame({
    'Actual': y_test.values,
    'LR_Pred': lr_preds,
    'LR_Prob': lr_probs,
    'RF_Pred': rf_preds,
    'RF_Prob': rf_probs,
    'DNN_Pred': dnn_preds,
    'DNN_Prob': dnn_probs
})

# View first 20 predictions
print("Sample Predictions:")
print("=" * 80)
print(predictions_df.head(20).to_string(index=False))

```

Sample Predictions:

Actual	LR_Pred	LR_Prob	RF_Pred	RF_Prob	DNN_Pred	DNN_Prob
1	1	0.558936	1	0.604994	0	0.415715
0	0	0.479679	1	0.521751	1	0.671899
0	0	0.310861	0	0.441351	0	0.286902
1	1	0.507620	1	0.657821	1	0.585613
0	1	0.586324	1	0.520562	1	0.679461
1	0	0.483201	0	0.308577	1	0.501587
0	0	0.395235	0	0.197001	0	0.222004
1	0	0.421257	1	0.555085	0	0.343363
0	1	0.513560	1	0.509744	1	0.532644
0	0	0.300240	0	0.399589	0	0.288548
0	1	0.599736	1	0.599173	1	0.685539
0	0	0.397359	0	0.439991	0	0.374813
1	0	0.492876	1	0.502165	1	0.562245
1	0	0.422962	1	0.675033	1	0.597472
1	0	0.349140	0	0.387403	0	0.229550
1	1	0.682840	0	0.382601	1	0.680490
0	0	0.289784	0	0.375985	0	0.245283
0	1	0.664613	1	0.535018	1	0.596997
1	0	0.437240	0	0.410950	0	0.483759
0	0	0.471336	1	0.613439	1	0.532856

```

In [14]: # 6.3. Cases where models agree or differ
print("\n" + "="*80)
print("DETAILED PREDICTION ANALYSIS")
print("="*80)

# Case 1: All models agree - correct
correct_agreement = predictions_df[
    (predictions_df['Actual'] == predictions_df['LR_Pred']) &
    (predictions_df['LR_Pred'] == predictions_df['RF_Pred']) &
    (predictions_df['RF_Pred'] == predictions_df['DNN_Pred'])
]

```

```

print(f"\n All models agree AND correct: {len(correct_agreement)} cases")
print(correct_agreement.head(3))

# Case 2: All models agree - wrong
wrong_agreement = predictions_df[
    (predictions_df['Actual'] != predictions_df['LR_Pred']) &
    (predictions_df['LR_Pred'] == predictions_df['RF_Pred']) &
    (predictions_df['RF_Pred'] == predictions_df['DNN_Pred'])
]
print(f"\n All models agree BUT wrong: {len(wrong_agreement)} cases")
print(wrong_agreement.head(3))

# Case 3: DNN correct when others wrong
dnn_saves = predictions_df[
    (predictions_df['Actual'] == predictions_df['DNN_Pred']) &
    (predictions_df['LR_Pred'] != predictions_df['Actual']) &
    (predictions_df['RF_Pred'] != predictions_df['Actual'])
]
print(f"\n DNN correct when LR & RF wrong: {len(dnn_saves)} cases")
print(dnn_saves.head(3))

```

=====
DETAILED PREDICTION ANALYSIS
=====

All models agree AND correct: 8336 cases

	Actual	LR_Pred	LR_Prob	RF_Pred	RF_Prob	DNN_Pred	DNN_Prob	
2	0	0	0.310861		0	0.441351	0	0.286902
3	1	1	0.507620		1	0.657821	1	0.585613
6	0	0	0.395235		0	0.197001	0	0.222004

All models agree BUT wrong: 4240 cases

	Actual	LR_Pred	LR_Prob	RF_Pred	RF_Prob	DNN_Pred	DNN_Prob	
4	0	1	0.586324		1	0.520562	1	0.679461
8	0	1	0.513560		1	0.509744	1	0.532644
10	0	1	0.599736		1	0.599173	1	0.685539

DNN correct when LR & RF wrong: 802 cases

	Actual	LR_Pred	LR_Prob	RF_Pred	RF_Prob	DNN_Pred	DNN_Prob	
5	1	0	0.483201		0	0.308577	1	0.501587
22	0	1	0.513963		1	0.574623	0	0.481774
66	1	0	0.465984		0	0.495227	1	0.509838

In [15]: # 6.4. Compare training times
import time

```

# Train baseline
start = time.time()
lr.fit(X_train, y_train)
lr_time = time.time() - start

start = time.time()
rf.fit(X_train, y_train)
rf_time = time.time() - start

# Train DNN

```

```
start = time.time()
model.fit(X_train, y_train, epochs=20, verbose=0)
dnn_time = time.time() - start

print(f"Logistic Regression: {lr_time:.2f}s")
print(f"Random Forest: {rf_time:.2f}s")
print(f"DNN: {dnn_time:.2f}s ({dnn_time/lr_time:.1f}x slower than LR, {dnn_time/rf_time:.1f}x slower than RF)
print(f"\nPerformance gain - DNN vs. Logistic Regression: {(auc - lr_auc_score(y_test)):.2%} improvement
print(f"\nPerformance gain - DNN vs. Random Forest: {(auc - rf_auc_score(y_test)):.2%} improvement
```

Logistic Regression: 8.15s

Random Forest: 3.97s

DNN: 207.72s (25.5x slower than LR, 52.3x slower than RF)

Performance gain - DNN vs. Logistic Regression: 5.8% improvement

Performance gain - DNN vs. Random Forest: 3.3% improvement